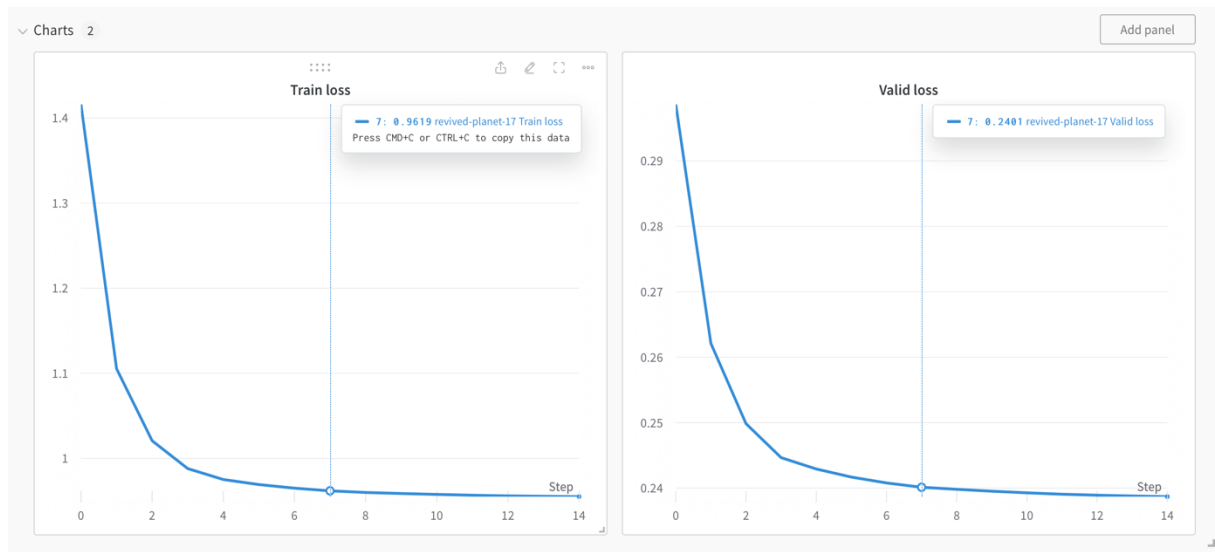


ASSIGNMENT 3 - DEEP LEARNING REPORT

- The transformation techniques I used:
 - `Resize(256, 256, interpolation=cv.INTER_LINEAR)`,
 - `HorizontalFlip(p=0.3)`
 - `VerticalFlip(p=0.3)`
 - `CoarseDropout(p=0.2, max_height=35, max_width=35, fill_value=255)`
 - `RandomSnow(snow_point_lower=0.1, snow_point_upper=0.15, brightness_coeff=1.5, p=0.09)`, `RandomShadow(p=0.1)`,
 - `Normalize()`
 - `ToTensorV2()`.
- The model is a U-Net with a pretrained Resnet50 encoder. I hoped that the results would exceed 0.7 accuracy but it did not happen as I expected.
- Link to my github repo: <https://github.com/cristiano2003/DL-Assignment-3.git>
- Below is the plot of my result:



- Link to my model checkpoint:
<https://www.kaggle.com/anhdaotruong/unet-training>

Note: In the ipynb file, I have referenced the data processing from the repo:
<https://github.com/GivraInNguyen/BKAI-IGH-Neopolyp-Segmentation.git>